

**Logistic Regression** is used for classification tasks. The two classes are usually 0 and 1 (false and true) but there are more extensions if there are more than 2 classes. The regression model fits the data into a logistic (sigmoidal curve) instead of fitting it into a straight line. The logistic regression assumes that the relationship between the feature and the class variable is non-linear.

**Ridge Regression** is a regularized version of Linear Regression and is also known as the Tikhonov Regularization. It uses the same equation as the Linear Regression Model to make predictions however the regression coefficients are chosen in such a way that not only does it fit the training data well, but it also makes sure that the magnitude of coefficients is as small as possible. Small values of coefficients means that each feature will have little effect on the outcome and the regression line will have a small slope. This makes it in such a way that the linear regression model is less complex and more restricted thus making it less likely to overfit. It uses the so called L2 regularization method. The main goal of this algorithm is to have high accuracy on the training data and be a low complexity model. The parameter  $\alpha$  controls the trade-off between the performance on the training set and the model complexity.

**Lasso Regression** is another type of regularized version of the standard Linear Regression Model. LASSO stands for Least Absolute Shrinkage and Selection Operator Regression. As the Ridge regression, this regression model adds a regularization term to the cost function, but it uses the L1 norm of the regression coefficient vector unlike the Ridge regression which uses L2 norm. The main goal of this regression model is to get high accuracy on training data while have a low complexity model. The consequence of using the L1 norm is that some regression coefficients can become exactly 0. That means that certain features would be completely ignored by the model, which makes it kind of like an automatic feature selection model. As the features would be less, the model itself will be less complex.

**Naive Bayes Algorithm** uses the Bayes theorem to solve classification tasks. It takes into consideration 2 assumptions, these are: The attributes are conditionally independent of each other for each class value, all attributes are equally important. Considering that these assumptions are almost never correct, the algorithm is called Naive Bayes theorem. These assumptions, however, leads to a simple and easy to implement algorithm which works well in practice. Given below are the steps to use Bayes Theorem: Calculate  $P(H|E)$  for each  $H$  (class) -  $P(Yes|E)$  and  $P(no|E)$ . Compare them and assign  $E$  to the class with the highest probability. For  $P(H|E)$  we need to calculate  $P(H)$  and  $P(E|H)$  from the given training data.

**Evaluating Machine Learning Algorithms:** Holdout Method, Holdout Method with parameter tuning, Holdout Method with stratification, Repeated Holdout Method, Cross Validation, Cross Validation with stratification, Leave-one-out Validation, Cross validation for parameter tuning.

**Entropy** measures the homogeneity of a set with respect to the class. The smaller the entropy, the greater the purity of the data set. **Information Gain** measures the reduction in entropy caused by using an attribute to partition the set of training examples. The best attribute is the one with the highest information gain (biggest reduction in entropy). Information gain is the difference between 2 values of entropies.

**Ensemble Methods: Bagging** (also called as bootstrap aggregation. Given a dataset 'D' with 'n' examples. Bootstrap examples 'D-' contains 'n' examples as well chosen randomly from 'D' where some examples will appear more than once or some not at all.). **Ada Boosting** (This algorithm uses a weighted training set meaning that each training example has an associated weight to it which is greater or equal to 0. The higher the weight, the more difficult the example was to classify by the previous classifier and hence have a higher chance to be selected for the training set of the next classifier).

**Kernel trick** is a method for computing the dot product of a pair of vectors in the new space without first computing the transformation of each vector from the original to the new space. First, we compute the dot product of the original features and use it in a 'kernel function' to determine the dot product of the transformed features. The kernel function specifies the relationship between the dot products in the original and transformed space.

**Principal Component Analysis** is the most popular dimensionality reduction method. It is often called a feature projection method. The main idea of this method is to find a new set of dimensions and project data into it. The dimensionality of the new space is smaller than the original space and the new axes captures the essence of the data. The resulting dataset (projection) can be used as an input to train a ML algorithm. In short, the PCA method deals with the construction of new features which are smaller than the number of original features.

**Back propagation Algorithm** is used to train a multi-layer perceptron neural network. The main idea of this algorithm is that for each training algorithm, an input is propagated through the network and an output is calculated. This is compared with the target output and error is calculated. The weights are then updated to reduce the error until the error for overall examples is less than the threshold value. The weights are updated backwards, from the output to the input neurons by propagating the weight change to minimize the error and hence is called a back propagation algorithm.

**Vanishing Gradient Problem:** when there are more layers in the network, the value of the product of derivative decreases until at some point the partial derivative of the loss function approaches a value close to zero, and the partial derivative vanishes. We call this the vanishing gradient problem.

**Dropout:** This is a method to prevent overfitting. The main idea is to avoid learning spurious features at the hidden nodes through intuition. Relevant features are more resilient to the removal of neurons, and they perform well for different combinations of neurons. While spurious features depend on certain neurons. Dropout forces the Neural Network to be less dependent on certain neurons to collect more evidence from the other neurons and to be more robust to noise. During training, at each iteration of the backpropagation, random neurons are selected in each layer and their values are set to 0. This results in a thinned sub-network of a smaller size. During training, the weights and biases are updated using back propagation algorithm and then new weight values are added to the original network. During testing, no neurons are dropped out and the network is scaled down based on the dropout rate.

**SoftMax:** The neural network outputs are processed to turn them into probabilities. The main motivation for this is to interpret the outputs as probabilities that sum up to 1.

**GMM** stands for **Gaussian Mixture Model** clustering which is a probabilistic clustering technique. It assumes that the data is generated by a mixture of normal distributions. In this algorithm we assume that the data is generated by a mixture of  $k$  Gaussian distributions and each distribution has 2 parameters, mean and standard deviation. One distribution corresponds to one cluster and starting from the initial values, the parameters are estimated iteratively. After each estimation, probability for each example is computed and then parameters are recomputed until they don't change. K-Means uses hard assignment whereas GMM uses probabilistic assignment. GMM can be seen as a generalization of K-Means. It is more flexible and it allows for elliptical clusters rather than circular for probabilistic assignment to each cluster rather than crisp.

**Agglomerative clustering (Hierarchical)** also called as the bottom-up clustering wherein the clusters merges iteratively. The algorithm starts with each item in its own cluster and iteratively merge until all items belong in one big cluster. The key operation in this type of clustering algorithm is computing the distance between 2 clusters. There are different versions of how the clusters are merged at each step.

**Divisive clustering (Hierarchical)** also called the top-down clustering wherein the clusters split into two until all items are in their own clusters. This can be implemented based on computing the minimum spanning tree.

**DBSCAN** stands for Density Based Spatial Clustering of Applications with Noise. The clusters created are regions of high density, separated from one another by regions with low density. In contrast to K-Means, DBSCAN can find clusters with arbitrary and complex shape. The main idea of the algorithm was that in a cluster the density of the points around it should be higher than the threshold. We need to define the density and the neighbourhood of a point. The neighbourhood of a point is the area within a radius of the point. Density of a point is the number of points in the neighbourhood of a point including that point. And Density threshold MinPts is the minimum number of points in the EPS neighbourhood of a point. The density of a point depends on the neighbourhood Eps (Epsilon). If the epsilon is too big, the density will be less and if the Eps is too small, all the points will have density as 1.

**Grid based clustering** is a density-based clustering algorithm which utilizes grids. The main idea is to break the data points into grid cells and then form clusters from the cells that are dense enough. The values of each attribute are split into intervals which are called grid cells. The common approach to do this is by splitting the values in equal width intervals. This leads the grid cells to have the same volume and the number of points will give the density of the cell. There are more sophisticated approaches as well such as breaking the values of an attribute in intervals such that there are equal data points in each grid, using clustering to determine intervals, break the initial values into large number of equal width intervals and then combine intervals with similar densities.

Clustering algorithms that use a few attributes from a pool of attributes is called **subspace clustering algorithms**. **CLAQUE** is called as the Clustering in Quest algorithm. It is a dimension growth-based clustering algorithm. The main idea of the algorithm is to find subspaces of high dimensionality where high density clusters exist. At each subspace, the data is partitioned into rectangular cells and the dense cells are identified based on a threshold. To efficiently find the dense cells at  $k$ -dim subspace, it uses the Apriori property from association rule mining and only considers the dense cells from the previous lower dimensional subspace ( $k-1$ ). Apriori principle states that if a set of points forms a density-based cluster in  $k$  dimensional space, then the same set of points are part of the density-based cluster in all possible subsets of those 3 dimensions.

**Markov Model** was proposed by Andrey Markov. A Markov chain (Markov process) is a model describing a sequence of transitions from one state to another, in which the probability of each state depends on the previous state. The main assumption of Markov model is that the probability of any concrete state depends only on the previous state and not the older history.

**Hidden Markov Models:** Markov Models are useful when probability of directly observable states is to be computed. Hidden Markov Models are used when are not observed directly or are hidden but can develop a judgement based on indirect observations. A Hidden Markov model is a probabilistic model that allow us to predict a sequence of hidden events from a set of observed events.

**Markov's Decision Process** is a discrete-time stochastic control process. It provides a mathematical framework for modelling decision making in situations where outcomes are partly random and partly under the control of a decision maker.


#### Question 5.1R

Given the training data in the table below where *credit history*, *debt*, *deposit* and *income* are attributes and *risk* is the class, predict the class of the following new example using the 1R algorithm: *credit history=unknown*, *debt=low*, *deposit=none*, *income=average*. If needed, settle ties by random selection. Show your calculations.

credit history	debt	deposit	income	risk
bad	high	none	low	high
unknown	high	none	average	high
unknown	low	none	average	moderate
unknown	low	none	low	high
unknown	low	none	high	low
unknown	low	adequate	high	low
bad	low	none	low	high
bad	low	adequate	high	moderate
good	low	none	high	low
good	high	adequate	high	low
good	high	none	low	high
good	high	none	average	moderate
good	high	none	high	low
bad	high	none	average	high

#### Solution:

- Attribute *credit history*  
bad: 0 low, 1 moderate, 3 high => risk=high, errors: 1/4  
unknown: 2 low, 1 moderate, 2 high => risk=low, errors: 3/5  
good: 3 low, 1 moderate, 1 high => risk=low, errors: 2/5  
total errors: 6/14
- Attribute *debt*  
high: 2 low, 1 moderate, 4 high => risk=high, errors: 3/7  
low: 3 low, 2 moderate, 2 high => risk=low, errors: 4/7  
total errors: 7/14
- Attribute *deposit*  
none: 3 low, 2 moderate, 6 high => risk=high, errors: 5/11  
adequate: 2 low, 1 moderate, 0 high => risk=low, errors: 1/3  
total errors: 6/14
- Attribute *income*




input vector x	initial weights w	initial biases θ
$x_1$	$w_{11}$ $w_{12}$ $w_{13}$ $w_{14}$ $w_{15}$ $w_{16}$ $w_{17}$ $w_{18}$ $w_{19}$ $w_{10}$	$\theta_1$ $\theta_2$ $\theta_3$
$x_2$	$w_{21}$ $w_{22}$ $w_{23}$ $w_{24}$ $w_{25}$ $w_{26}$ $w_{27}$ $w_{28}$ $w_{29}$ $w_{20}$	$\theta_1$ $\theta_2$ $\theta_3$
$x_3$	$w_{31}$ $w_{32}$ $w_{33}$ $w_{34}$ $w_{35}$ $w_{36}$ $w_{37}$ $w_{38}$ $w_{39}$ $w_{30}$	$\theta_1$ $\theta_2$ $\theta_3$

Input to neuron 4:  $z_4 = 0.2 \times 0.4 + 1 \times (-0.5) + 0.4 = -0.7$ , output of neuron 4:  $o_4 = 1/(1 + e^{0.7}) = 0.332$   
 Input to neuron 5:  $z_5 = 1 \times (-0.3) + 0.9 \times 0.1 + 0.2 \times 0.2 = 0.1$ , output of neuron 5:  $o_5 = 1/(1 + e^{0.1}) = 0.525$  **NN output**  
 Input to neuron 6:  $z_6 = 0.332 \times (-0.3) + 0.525 \times (-0.2) + 0.1 = -0.105$ , output of neuron 6:  $o_6 = 1/(1 + e^{0.105}) = 0.474$

$\delta_6 = (t_6 - o_6) \times o_6 \times (1 - o_6) = (1 - 0.474) \times 0.474 \times (1 - 0.474) = 0.1311$   
 $\Delta w_{64} = \eta \times \delta_6 \times o_4 = 0.9 \times 0.1311 \times 0.332 = 0.039$ ,  $w_{64} \text{ new} = w_{64} \text{ old} + \Delta w_{64} = -0.3 + 0.039 = -0.261$   
 $\Delta w_{65} = \eta \times \delta_6 \times o_5 = 0.9 \times 0.1311 \times 0.525 = 0.0619$ ,  $w_{65} \text{ new} = w_{65} \text{ old} + \Delta w_{65} = -0.2 + 0.0619 = -0.138$   
 $\theta_6 \text{ new} = \theta_6 \text{ old} + \Delta \theta_6 = \theta_6 \text{ old} + \eta \times \delta_6 \times 1 = 0.1 + 0.9 \times 0.1311 \times 1 = 0.218$   
 $\delta_5 = o_5 \times (1 - o_5) \times w_{56} = 0.332 \times (1 - 0.332) \times (-0.3) \times 0.1311 = -0.0087$   
 $\Delta w_{54} = \eta \times \delta_5 \times o_4 = 0.9 \times (-0.0087) \times 1 = -0.0079$ ,  $w_{54} \text{ new} = w_{54} \text{ old} + \Delta w_{54} = 0.2 - 0.0079 = 0.1921$   
 $\Delta w_{55} = \eta \times \delta_5 \times o_5 = 0.9 \times (-0.0087) \times 0 = 0$ ,  $w_{55} \text{ new} = w_{55} \text{ old} + \Delta w_{55} = 0.4 + 0 = 0.4$   
 $\Delta w_{46} = \eta \times \delta_4 \times o_6 = 0.9 \times (-0.0087) \times 1 = -0.0079$ ,  $w_{46} \text{ new} = w_{46} \text{ old} + \Delta w_{46} = -0.5 - 0.0079 = -0.5079$   
 $\theta_4 \text{ new} = \theta_4 \text{ old} + \Delta \theta_4 = \theta_4 \text{ old} + \eta \times \delta_4 \times 1 = 0.4 + 0.9 \times (-0.0087) \times 1 = 0.4078$

#### Question 8. Markov models

Given is the following Markov model for the weather in Sydney:



a) Given that today the weather is *Sunny*, what is the probability that it will be *Sunny* tomorrow and *Rainy* the day after tomorrow, i.e. what is the probability  $P(\pi_2 = \text{Rainy}, \pi_3 = \text{Sunny} | \pi_1 = \text{Sunny})$ ?

Hint:  $P(A,B|C) = P(A|B,C) P(B|C)$

b) If the weather yesterday was *Rainy*, and today is *Foggy*, what is the probability that tomorrow it will be *Sunny*?

For both questions, briefly show your calculations.

#### Solution:

a)  
 $P(\pi_2 = \text{Rainy}, \pi_3 = \text{Sunny} | \pi_1 = \text{Sunny}) =$   
 $= P(\pi_2 = \text{Rainy} | \pi_2 = \text{Sunny}, \pi_1 = \text{Sunny}) * P(\pi_2 = \text{Sunny} | \pi_1 = \text{Sunny}) = (\text{rule from hint})$   
 $= P(\pi_2 = \text{Rainy} | \pi_2 = \text{Sunny}) * P(\pi_2 = \text{Sunny} | q, \pi_1 = \text{Sunny}) = (\text{Markov assumption})$   
 $= 0.05 * 0.8 = 0.04$

b)  
 $P(\pi_3 = \text{Sunny} | \pi_2 = \text{Foggy}, \pi_1 = \text{Rainy}) =$   
 $= P(\pi_3 = \text{Sunny} | \pi_2 = \text{Foggy}) = (\text{Markov assumption})$   
 $= 0.2$

### Exercise 1: Nearest Neighbor (to do in class)

The dataset below consists of 4 examples described with 3 numeric features (a1, a2 and a3); the class has 2 values: yes and no.

What will be the prediction of 1-Nearest Neighbor (1-NN) and 3-Nearest Neighbor (3-NN) with Euclidean distance for the following new example: a1=2, a2=4, a3=2?

temperature=1.0  
humidity=0.2

Assume that all attributes are measured on the same scale – no need for normalization.

a1	a2	a3	class
1	1	3	yes
2	1	5	yes
3	2	2	no
4	3	1	no

Exercise adapted from M. Kolar, Introduction to Machine Learning, Springer, 2017

### Solution:

Distance:  $\sqrt{(2-1)^2 + (4-1)^2 + (2-3)^2} = \sqrt{5}$   
Distance:  $\sqrt{(2-2)^2 + (4-5)^2 + (2-2)^2} = \sqrt{2}$   
Distance:  $\sqrt{(2-3)^2 + (4-2)^2 + (2-1)^2} = \sqrt{5}$   
Distance:  $\sqrt{(2-4)^2 + (4-3)^2 + (2-1)^2} = \sqrt{5}$

The closest nearest neighbor is ex. 2, hence 1-NN predicts class=yes  
The closest 3 nearest neighbors are ex.2 (yes), ex.1 (yes) and ex.3 (no); the majority class is yes.  
Hence, 3-NN predicts class= yes

### Exercise 2: Naïve Bayes with nominal features (to do at your own time)

Consider the iPhone dataset given below. There are 4 nominal attributes (age, income, student, and credit\_rating) and the class is buys\_iPhone with 2 values: yes and no.

What would be the prediction of 1-NN and 3-NN for the following new example:  
age<30, income=medium, student=yes, credit\_rating=fair

If there are ties, make random selection.

Tip: As the examples are described with nominal attributes, when calculating the distance use the following rule:

distance=1 if the attributes differ by 2 values that are not the same  
distance=0 if the attributes differ by 1 value that are the same  
e.g. D1, new=0+1+1+0=2 (age=2)

age	income	student	credit_rating	buys_iPhone
<30	high	no	excellent	no
<30	high	no	excellent	no
<30	high	no	fair	yes
<30	medium	no	fair	yes
<30	medium	no	good	yes
<30	low	yes	excellent	yes
<30	low	yes	good	yes
<30	low	yes	fair	no
<30	low	yes	excellent	no
>30	medium	no	excellent	no

Dataset adapted from J. Han and M. Kamber, Data Mining, Concepts and Techniques, Morgan Kaufmann.

### Solution:

new example: age<30, income=medium, student=yes, credit\_rating=fair

D1, new = 0+1+1+1+0 = 3 (age=2)  
D2, new = 0+1+1+1+1 = 4 (age=3)  
D3, new = 0+1+1+1+0 = 3 (age=2)  
D4, new = 0+1+1+1+0 = 3 (age=2)  
D5, new = 0+1+1+1+1 = 4 (age=3)  
D6, new = 0+1+1+1+1 = 4 (age=3)  
D7, new = 0+0+0+0+0 = 0 (age=1)  
D8, new = 0+1+1+1+1 = 4 (age=3)  
D9, new = 0+1+1+1+1 = 4 (age=3)

- 1-NN: ex. 7 (D1) is the closest neighbor, hence 1-NN predicts buy\_iPhone= yes
- 3-NN: the 3 closest neighbors are: ex. 7 (D1), ex. 1 and ex. 4 (D=0+2+1)=2, no and 1 yes => the majority class is no. Hence, 3-NN predicts buy\_iPhone= no

### Exercise 3: PRISM for the at your own time

Given the training data in the table below, generate the PRISM rules for class=no. In case of ties, make random selection.

Weather data with nominal attributes:

	outlook	temperature	humidity	windy	play
1.	sunny	hot	high	false	no
2.	sunny	hot	high	true	no
3.	overcast	hot	high	false	yes
4.	rainy	mild	high	false	yes
5.	rainy	cool	normal	false	yes
6.	rainy	cool	normal	true	no
7.	overcast	cool	normal	false	yes
8.	overcast	cool	normal	true	yes
9.	sunny	mild	normal	false	yes
10.	sunny	mild	normal	false	yes
11.	sunny	mild	normal	true	yes
12.	overcast	mild	high	false	yes
13.	overcast	hot	normal	false	yes
14.	rainy	mild	high	true	no

### Solution:

Let's start with generating rules for class no:

if ? then class=no

10 possible tests with their corresponding accuracy p:

outlook=sunny 3/5  
outlook=overcast 0/4  
outlook=rainy 2/5

temperature=hot 2/4  
temperature=cool 2/5  
temperature=mild 1/5

humidity=high 4/7  
humidity=normal 1/7

windy=true 3/6  
windy=false 2/8

Best test: outlook=sunny

Rule: if outlook=sunny then class=no

Examples covered by the current rule:

outlook temperature humidity windy play  
1. sunny hot high false no  
2. sunny hot high true no  
8. sunny cool normal false yes  
9. sunny cool normal true no  
11. sunny mild normal true yes

Not a perfect rule as it covers also 2 examples from class yes => add other tests to this rule if outlook=sunny and 7 then class=no

Possible tests with the corresponding accuracy p:

temperature=hot 2/2  
temperature=cool 1/1  
temperature=mild 0/2

humidity=high 3/3  
humidity=normal 0/2

windy=true 2/2  
windy=false 2/3

Best test: humidity=high (bigger coverage than temperature=hot and temperature=cool)

Rule: if outlook=sunny and humidity=high then class=no, prefer rule as it covers only examples from class no => stop adding other tests to this rule, delete the examples covered by the rule, there are still uncovered examples from class no, so generate another rule for class no.

outlook temperature humidity windy play  
3. overcast hot high false yes  
4. rainy mild high false yes  
6. rainy cool normal true no  
7. overcast cool normal false yes  
10. sunny mild normal false yes  
12. overcast mild high false yes  
13. overcast hot normal false yes  
14. rainy mild high true no

if ? then class=no

10 possible tests with their corresponding accuracy p:

outlook=overcast 0/2  
outlook=rainy 2/5  
outlook=sunny 0/2

temperature=hot 0/2  
temperature=cool 1/4  
temperature=mild 1/5

humidity=high 1/4  
humidity=normal 1/7

windy=true 2/5  
windy=false 0/6

Best test: outlook=overcast (draw with windy=true, random selection)

Rule: if outlook=overcast then class=no

Examples covered by the current rule:

outlook temperature humidity windy play  
3. overcast hot high false yes  
4. rainy mild high false yes

6. sunny cool normal true yes  
10. sunny cool normal true yes  
14. rainy mild high true no

Not a perfect rule as it covers also 3 examples from class yes => add other tests to this rule if outlook=overcast and 7 then class=no

Possible tests with the corresponding accuracy p:

temperature=hot 1/3  
temperature=cool 1/2

humidity=high 1/2  
humidity=normal 1/3

windy=true 2/2  
windy=false 0/3

Best test: windy=true

Rule: if outlook=overcast and windy=true then class=no, prefer rule as it covers only examples from class no => stop adding other tests to this rule, delete the examples covered by the rule, there are no more uncovered examples from class no, so the rule for class no are:

if outlook=overcast and windy=true then class=no  
if outlook=overcast and humidity=high then class=no

Repeat the procedure for class yes. PRISM will generate the following rules:

if outlook=overcast then class=yes  
if humidity=normal and windy=false then class=yes  
if temperature=mild and humidity=normal then class=yes  
if temperature=mild and windy=false then class=yes

### Exercise 4: Naïve Bayes for data with nominal features (to do in class)

Consider the following dataset where from=down is the class. Predict the class of the following new example using Naïve Bayes:

home owner = no, marital status = married, annual income=very high

Dataset adapted from, Tan, Karthikeyan and Kumar, Introduction to Data Mining, Pearson, 2019

Solution:  
E=house owner=no, marital status=married, annual income=very high  
E1 is house owner=no, E2 is marital status=married, E3 is annual income=very high

We need to compute P(E1|yes) and P(E2|no) and compare them.

$P(yes|E) = \frac{P(E_1|yes)P(E_2|yes)P(E_3|yes)}{P(E)}$   
 $P(no|E) = \frac{P(E_1|no)P(E_2|no)P(E_3|no)}{P(E)}$

P(yes)=5/10 P(no)=5/10

P(E1|yes)=P(house owner=yes)=3/5  
P(E2|yes)=P(marital status=married|yes)=1/5  
P(E3|yes)=P(annual income=very high|yes)=1/5

$P(yes|E) = \frac{3 \cdot 1 \cdot 1 \cdot 5}{5 \cdot 5 \cdot 5 \cdot 10} = \frac{3}{250} = \frac{P(E)}{P(E)}$

$P(no|E) = \frac{3 \cdot 3 \cdot 5 \cdot 9}{5 \cdot 5 \cdot 5 \cdot 10} = \frac{27}{250} = \frac{P(E)}{P(E)}$

P(no|E)=P(yes)=No: Naïve Bayes predicts loan default = no for the new example.

### Exercise 2: Naïve Bayes for data with numeric features (to do in class)

Consider the following dataset where from=down is the class. Predict the class of the following new example using Naïve Bayes:

home owner = no, marital status = married, annual income=120

Solution:  
1) Calculate the mean  $\mu$  and standard deviation  $\sigma$  values for the numeric feature income:

$\mu = \frac{\sum X_i}{n}$   $\sigma^2 = \frac{\sum (X_i - \mu)^2}{n-1}$

where  $X_i = 1, 2, \dots, n$  – the  $i$ -th measurement,  $n$  – number of measurements

We need to calculate the mean and standard deviation separately for each class (yes and no) – separate the values of income:

class yes income: 125, 70, 100, 120, 95, 60, 65, 75, 90, 75  
class no income: 120, 100, 120, 95, 60, 65, 75, 90, 75

$\mu_{income=yes} = 99$   $\mu_{income=no} = 109$   
 $\sigma_{income=yes} = 15.57$   $\sigma_{income=no} = 66.18$

2) Calculate P(income=120|yes) and P(income=120|no) using the probability density function for normal distribution:

$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$   
 $f(income=120|yes) = \frac{1}{15.57\sqrt{2\pi}} e^{-\frac{(120-99)^2}{2 \cdot 15.57^2}} = 0.01032$   
 $f(income=120|no) = \frac{1}{66.18\sqrt{2\pi}} e^{-\frac{(120-109)^2}{2 \cdot 66.18^2}} = 0.00995$

3) Calculating the probabilities P(yes|E) and P(no|E) using the Bayes Theorem, we already have the probabilities for the nominal attributes from the previous exercise:

$P(yes|E) = \frac{3}{5} \cdot \frac{0.01032}{5} \cdot \frac{5}{10} = \frac{0.00619}{10}$   
 $P(no|E) = \frac{3}{5} \cdot \frac{0.00995}{5} \cdot \frac{5}{10} = \frac{0.00619}{10}$

P(yes|E)=P(yes)=No: Naïve Bayes predicts loan default = no for the new example.

### Exercise 1: Decision trees and information gain (parts a and b) – done in class; the rest in your own time

Consider the following set of training examples:

shape color class  
circle blue +  
square blue +  
triangle blue -  
square red +  
square blue -  
square red +  
square blue -  
circle red +

What is the entropy of this collection of training examples with respect to the class?

Which attribute will be selected as root of the tree based on information gain?

Build the whole decision tree. Draw the tree after each selected attribute.

You may use this table to calculate information gain:

$K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$

1 2 0.50 4 5 0.26 6 7 0.19 8 9 0.47

1 3 0.58 1 6 0.43 1 3 0.38 7 9 0.28

2 3 0.39 5 6 0.22 3 8 0.53 8 9 0.15

1 4 0.5 1 7 0.40 5 8 0.42 1 10 0.33

3 4 0.31 2 7 0.52 7 9 0.17 8 10 0.52

1 5 0.46 3 7 0.42 1 9 0.35 7 10 0.36

2 5 0.53 4 7 0.48 2 9 0.48 9 10 0.14

3 1 0.58 4 7 0.55 4 9 0.52

Solution:  
a)  $H(S) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
b) Split on shape:  
 $H(S_{shape}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$   
 $H(S_{shape=blue}) = -[0.24 \log_2(0.24) + 0.76 \log_2(0.76)] = 0.9183$   
 $H(S_{shape=red}) = -[0.1 \log_2(0.1) + 0.9 \log_2(0.9)] = 0.469$   
gain(shape) = 0.9183 - 0.469 = 0.4493 bits

c) What is the entropy of this collection of training examples with respect to the class?

Which attribute will be selected as root of the tree based on information gain?

Build the whole decision tree. Draw the tree after each selected attribute.

You may use this table to calculate information gain:

$K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$   $K$   $Y$   $-log_2(Y)$

1 2 0.50 4 5 0.26 6 7 0.19 8 9 0.47

1 3 0.58 1 6 0.43 1 3 0.38 7 9 0.28

2 3 0.39 5 6 0.22 3 8 0.53 8 9 0.15

1 4 0.5 1 7 0.40 5 8 0.42 1 10 0.33

3 4 0.31 2 7 0.52 7 9 0.17 8 10 0.52

1 5 0.46 3 7 0.42 1 9 0.35 7 10 0.36

2 5 0.53 4 7 0.48 2 9 0.48 9 10 0.14

3 1 0.58 4 7 0.55 4 9 0.52

c) To answer this question we need to calculate the information gain of all attributes. The attribute with the highest information gain will be selected.

There are 2 attributes – shape and color. We already calculate the information gain for shape. Let's do this for color.

Split on color:  
 $H(S_{color}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
 $H(S_{color=blue}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$   
 $H(S_{color=red}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
gain(color) = 0.9183 - 0.5 = 0.4183 bits

gain(shape) > gain(color) => shape will be selected as the root of the DT (the first attribute to split on)

d) Building the decision tree:  
After selecting shape:  
circle square triangle  
1,2,8 3,5,6,7 4  
+ + + + + + + +  
needs further splitting

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log_2(0.75)] = 0.9183$   
For triangle:  $H(S_{triangle}) = -[0.33 \log_2(0.33) + 0.67 \log_2(0.67)] = 0.9183$

For circle:  $H(S_{circle}) = -[0.5 \log_2(0.5) + 0.5 \log_2(0.5)] = 1$   
For square:  $H(S_{square}) = -[0.25 \log_2(0.25) + 0.75 \log$